


# EfficientML.ai Lecture 03: Pruning and Sparsity

Part I



**Song Han**


Associate Professor, MIT  
Distinguished Scientist, NVIDIA

 @SongHan\_MIT






# Today's AI is too BIG!




Model Compression and Hardware Acceleration for Neural Networks: A Comprehensive Survey [Deng *et al.*, IEEE 2020]

# Efficient Deep Learning Techniques are Essential


Bridges the Gap between the Supply and Demand of Computation




# Part 1 of This Course: Efficient Inference



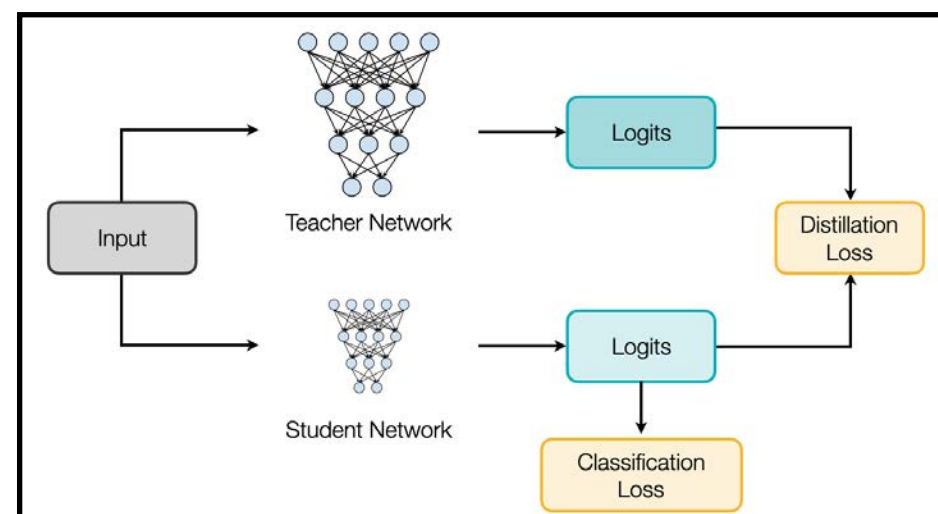
Pruning



Quantization



Neural Architecture Search



Knowledge Distillation

9/14 Lecture 3: **Pruning and Sparsity (Part I)**  
[ slides ] [ video ] [ video (live) ]

9/19 Lecture 4: **Pruning and Sparsity (Part II)**  
[ slides ] [ video ] [ video (live) ]

Lab 1 out

9/21 Lecture 5: **Quantization (Part I)**  
[ slides ] [ video ] [ video (live) ]

9/26 Lecture 6: **Quantization (Part II)**  
[ slides ] [ video ] [ video (live) ]

9/28 Lecture 7: **Neural Architecture Search (Part I)**  
[ slides ] [ video ] [ video (live) ]

Lab 1 due, Lab 2 out

10/3 Lecture 8: **Neural Architecture Search (Part II)**  
[ slides ] [ video ] [ video (live) ]

10/5 Lecture 9: **Knowledge Distillation**  
[ slides ] [ video ] [ video (live) ]

Lab 3 out

10/10 Student Holiday — No Class

10/12 Lecture 10: **MCUNet**  
[ slides ] [ video ] [ video (live) ]

Lab 2 due

10/17 Lecture 11: **TinyEngine**  
[ slides ] [ video ] [ video (live) ]

# MLPerf (the Olympic Game for AI Computing)

## Closed Division vs Open Division

- The open division submission on BERT: more than 4x while maintaining 99% accuracy.

	Closed Division	Open Division	Speedup
Offline samples/sec	1029	4609	4.5x

BERT Large performance metrics for both closed division and open division

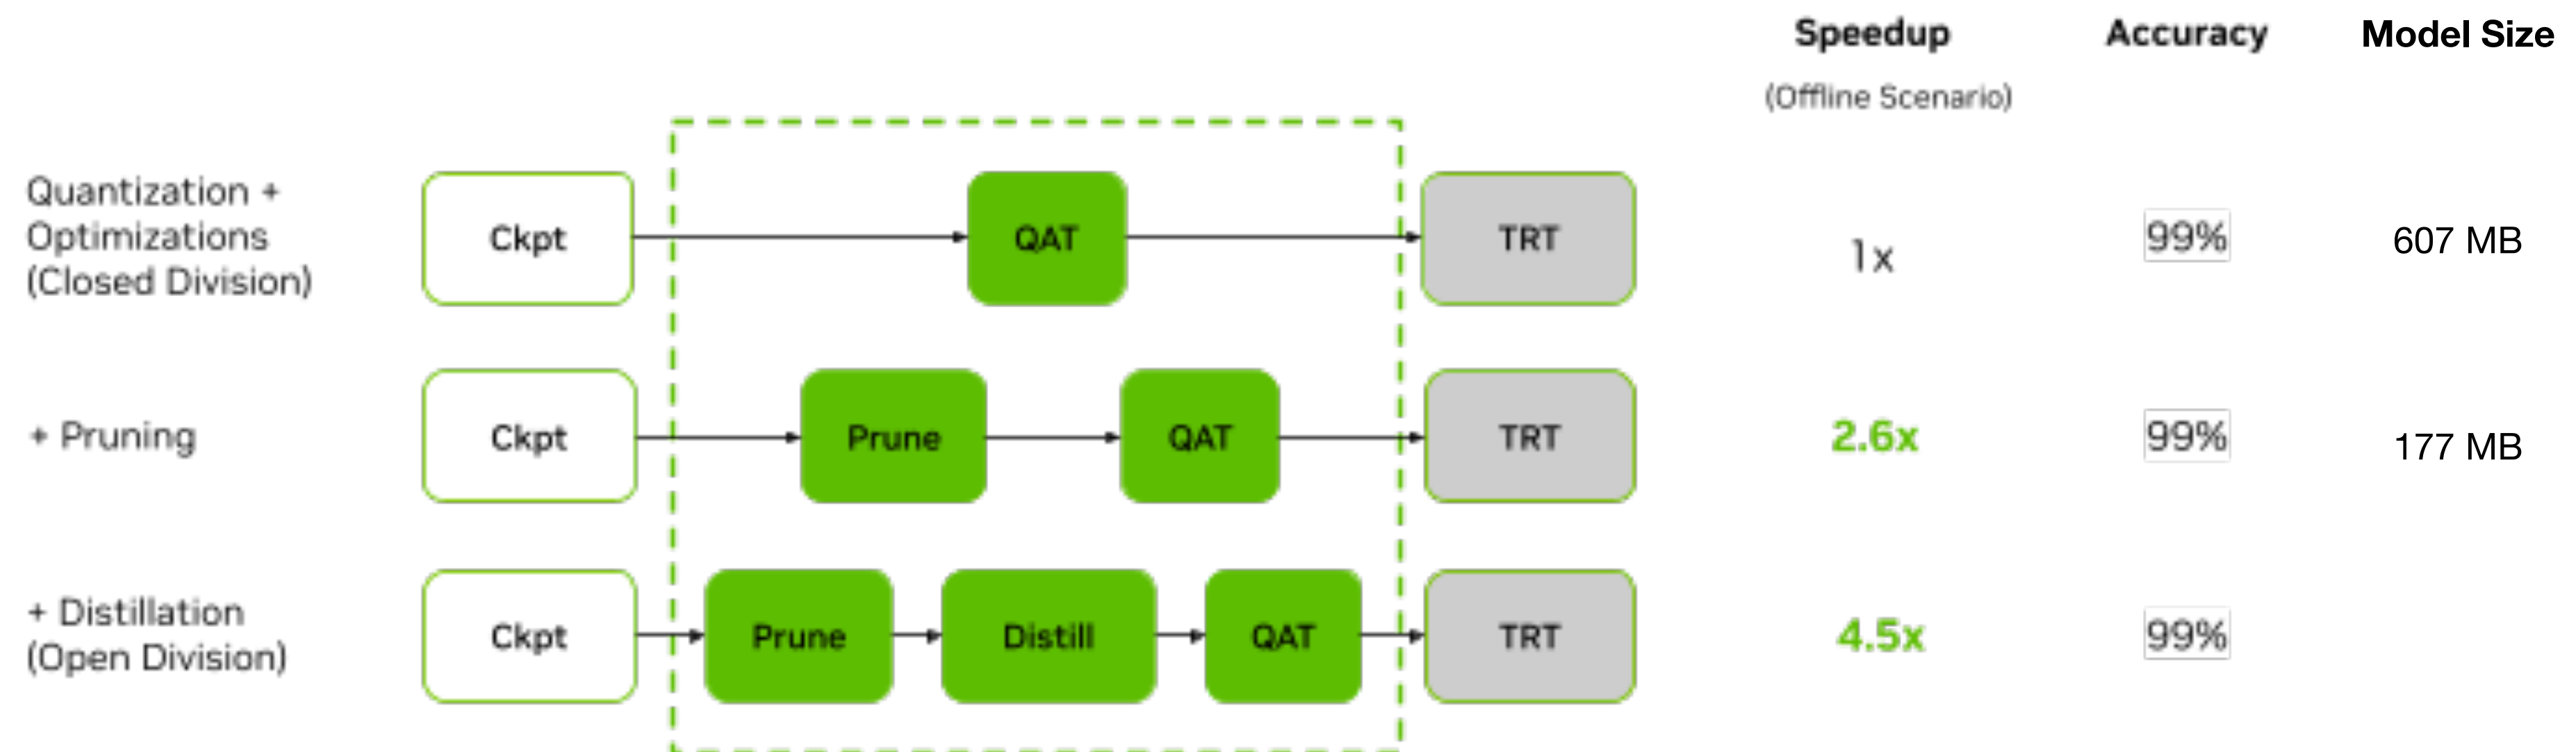
[Leading MLPerf Inference v3.1 Results with NVIDIA GH200 Grace Hopper Superchip Debut](#)



# MLPerf (the Olympic Game for AI Computing)

## Key techniques: pruning, distillation, quantization

- The open division submission on BERT: more than 4x while maintaining 99% accuracy.




Leading MLPerf Inference v3.1 Results with NVIDIA GH200 Grace Hopper Superchip Debut

# Lecture Plan


## Today we will:

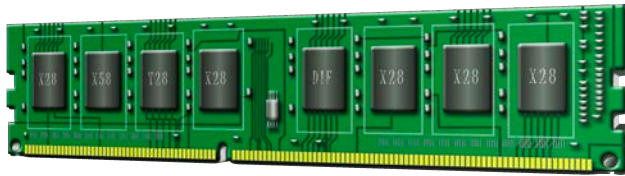
1. Introduce **neural network pruning** which can reduce the parameter counts of neural networks by more than 90%, decreasing the storage requirements and improving computation efficiency of neural networks.
2. Go through all steps of pruning, and introduce different **granularities** and **criteria** of neural network pruning.



# Memory is Expensive

Data Movement → More Memory Reference → More Energy



1  = 200 X +



This image is in the public domain

Computing's Energy Problem (and What We Can Do About it) [Horowitz, M., IEEE ISSCC 2014]



# Memory is Expensive

Data Movement → More Memory Reference → More Energy




Battery images are in the public domain  
[Image 1](#), [image 2](#), [image 2](#), [image 4](#)

Computing's Energy Problem (and What We Can Do About it) [Horowitz, M., IEEE ISSCC 2014]





# Neural Network Pruning

- **Introduction to Pruning**
  - What is pruning?
  - How should we formulate pruning?
- **Determine the Pruning Granularity**
  - In what pattern should we prune the neural network?
- **Determine the Pruning Criterion**
  - What synapses/neurons should we prune?
- **Determine the Pruning Ratio**
  - What should target sparsity be for each layer?
- **Fine-tune/Train Pruned Neural Network**
  - How should we improve performance of pruned models?



Learning Both Weights and Connections for Efficient Neural Network [Han *et al.*, NeurIPS 2015]

# Pruning Happens in Human Brain




Do We Have Brain to Spare? [Drachman DA, Neurology 2004]  
Peter Huttenlocher (1931–2013) [Walsh, C. A., Nature 2013]

Data Source: 1, 2  
Slide Inspiration: Alila Medical Media



# Neural Network Pruning

Make neural network smaller by removing synapses and neurons





Optimal Brain Damage [LeCun *et al.*, NeurIPS 1989]

Learning Both Weights and Connections for Efficient Neural Network [Han *et al.*, NeurIPS 2015]

# Neural Network Pruning



Make neural network smaller by removing synapses and neurons



Learning Both Weights and Connections for Efficient Neural Network [Han *et al.*, NeurIPS 2015]

# Neural Network Pruning

Make neural network smaller by removing synapses and neurons




Learning Both Weights and Connections for Efficient Neural Network [Han *et al.*, NeurIPS 2015]



# Neural Network Pruning


Make neural network smaller by removing synapses and neurons



Train Connectivity

Prune Connections



Train Weights



Learning Both Weights and Connections for Efficient Neural Network [Han et al., NeurIPS 2015]

# Neural Network Pruning

Make neural network smaller by removing synapses and neurons



Learning Both Weights and Connections for Efficient Neural Network [Han *et al.*, NeurIPS 2015]

# Neural Network Pruning


Make neural network smaller by removing synapses and neurons

Neural Network	#Parameters			MACs
	Before Pruning	After Pruning	Reduction	Reduction
AlexNet	61 M	6.7 M	9 ×	3 ×
VGG-16	138 M	10.3 M	12 ×	5 ×
GoogleNet	7 M	2.0 M	3.5 ×	5 ×
ResNet50	26 M	7.47 M	3.4 ×	6.3 ×
SqueezeNet	1 M	0.38 M	3.2 ×	3.5 ×



# Neural Network Pruning

Pruning the NeuralTalk LSTM does not hurt image caption quality.



**Baseline:** a basketball player in a white uniform is playing with a **ball**.

**Pruned 90%:** a basketball player in a white uniform is playing with a **basketball**.



**Baseline:** a brown dog is running through a grassy **field**.

**Pruned 90%:** a brown dog is running through a grassy **area**.



**Baseline:** a man **is riding a surfboard on a wave**.

**Pruned 90%:** a man **in a wetsuit is riding a wave on a beach**.




**Baseline:** a soccer player in red is running in the field.

**Pruned 95%:** a man **in a red shirt and black and white black shirt** is running through a field.



# Neural Network Pruning

Make neural network smaller by removing synapses and neurons



598 Le Cun, Denker and Solla

*Optimal Brain Damage*

Yann Le Cun, John S. Denker and Sara A. Solla  
AT&T Bell Laboratories, Holmdel, N. J. 07733

**Learning both Weights and Connections for Efficient Neural Networks**

Song Han  
Stanford University  
songhan@stanford.edu

Jeff Pool  
NVIDIA  
jpool@nvidia.com

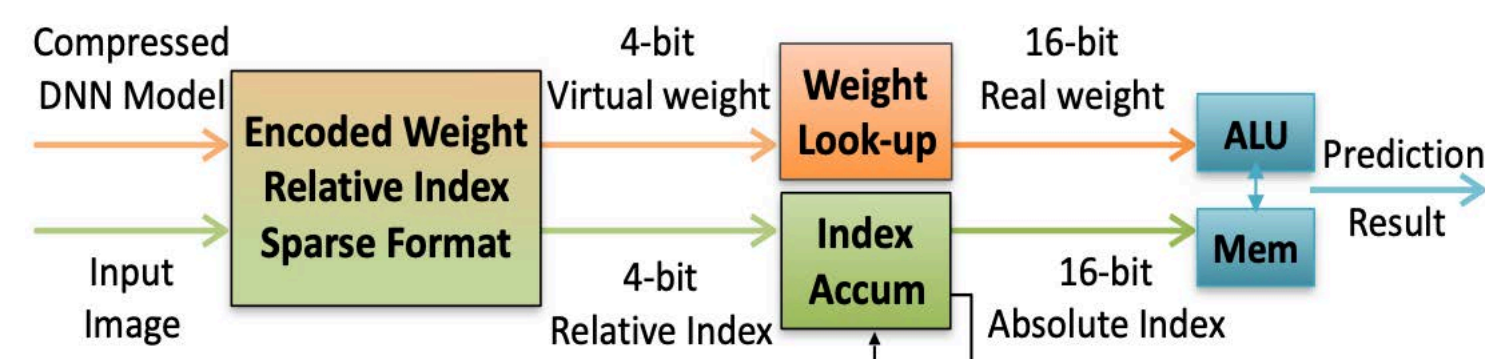
John Tran  
NVIDIA  
johntran@nvidia.com

William J. Dally  
Stanford University  
NVIDIA  
dally@stanford.edu

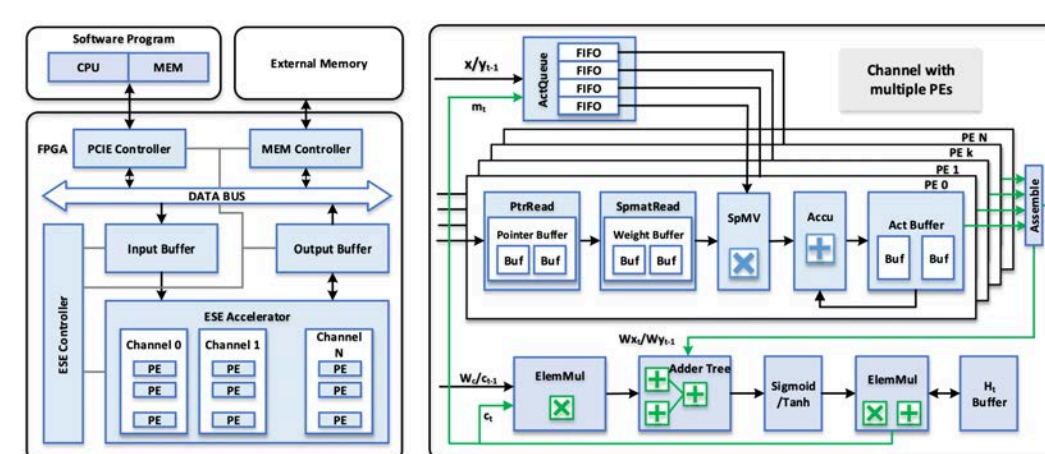
Source: <https://github.com/mit-han-lab/pruning-sparsity-publications>

# Pruning in the Industry


## Hardware support for sparsity




EIE [Han *et al.*, ISCA 2016]



ESE [Han *et al.*, FPGA 2017]



SpArch [Zhang *et al.*, HPCA 2020]  
SpAtten [Wang *et al.*, HPCA 2021]



2:4 sparsity in A100 GPU

2X peak performance, 1.5X measured BERT speedup

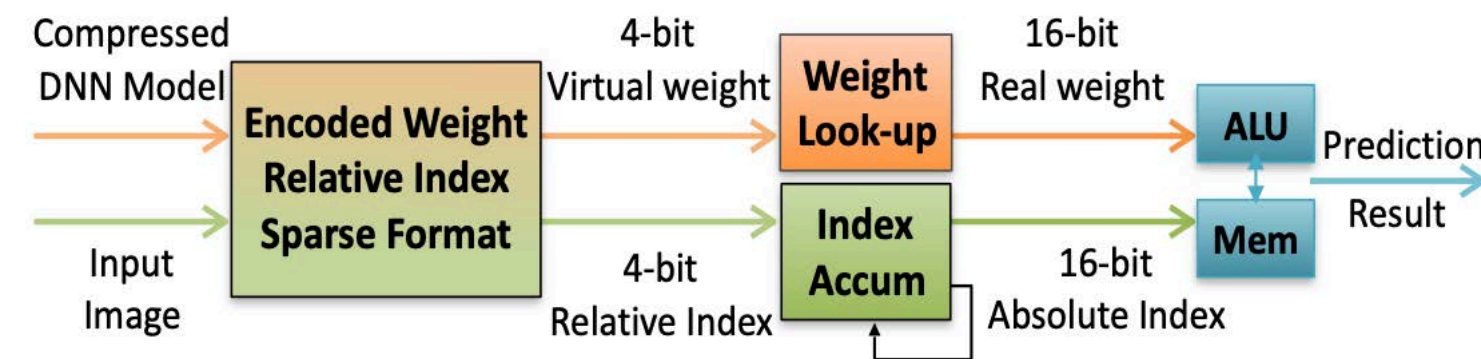


# Pruning in the Industry

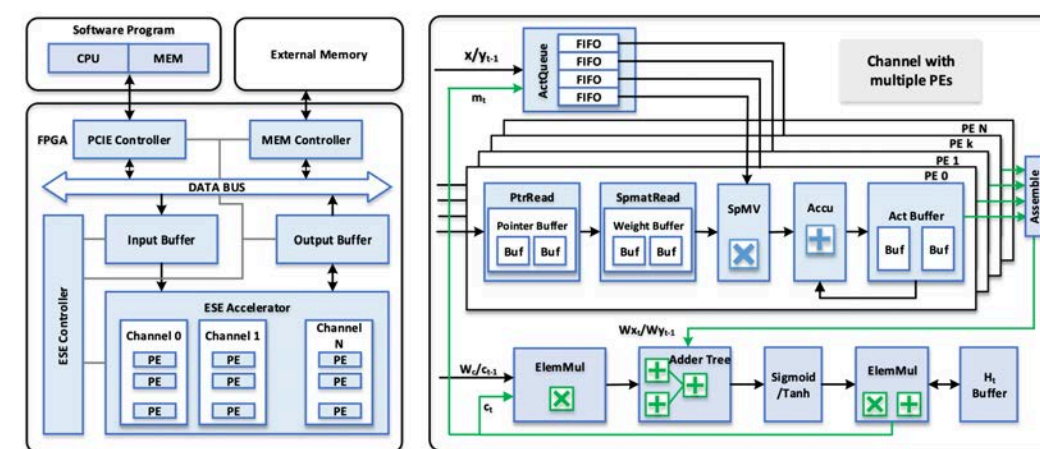
## Hardware support for sparsity

AMD  
XILINX


XILINX  
VITIS™




EIE [Han et al., ISCA 2016]




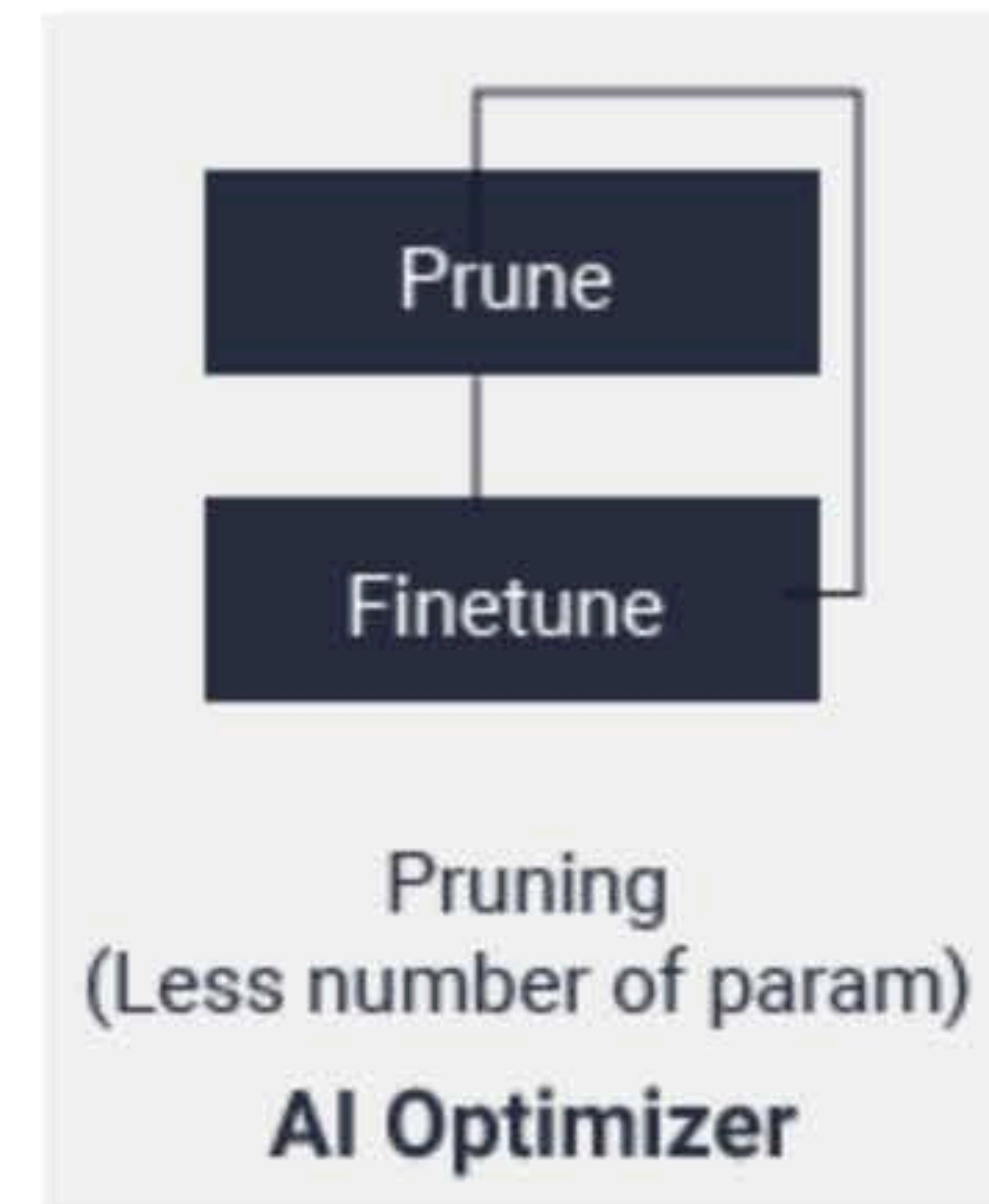
ESE [Han et al., FPGA 2017]



SpArch [Zhang et al., HPCA 2020]  
SpAtten [Wang et al., HPCA 2021]



Dense Neural Network (FP32)




Pruned Neural Network (FP32)

Reduce model complexity by 5x to 50x with minimal accuracy impact

# Neural Network Pruning

- **Introduction to Pruning**
  - What is pruning?
  - How should we formulate pruning?
- **Determine the Pruning Granularity**
  - In what pattern should we prune the neural network?
- **Determine the Pruning Criterion**
  - What synapses/neurons should we prune?
- **Determine the Pruning Ratio**
  - What should target sparsity be for each layer?
- **Fine-tune/Train Pruned Neural Network**
  - How should we improve performance of pruned models?



Learning Both Weights and Connections for Efficient Neural Network [Han *et al.*, NeurIPS 2015]

# Neural Network Pruning



- In general, we could formulate the pruning as follows:

$$\arg \min_{\mathbf{W}_P} L(\mathbf{x}; \mathbf{W}_P)$$

subject to

$$\|\mathbf{W}_P\|_0 < N$$


- $L$  represents the objective function for neural network training;
- $\mathbf{x}$  is input,  $\mathbf{W}$  is original weights,  $\mathbf{W}_P$  is pruned weights;
- $\|\mathbf{W}_P\|_0$  calculates the #nonzeros in  $\mathbf{W}_P$ , and  $N$  is the target #nonzeros.





# Neural Network Pruning


- **Introduction to Pruning**
  - What is pruning?
  - How should we formulate pruning?
- **Determine the Pruning Granularity**
  - In what pattern should we prune the neural network?
- **Determine the Pruning Criterion**
  - What synapses/neurons should we prune?
- **Determine the Pruning Ratio**
  - What should target sparsity be for each layer?
- **Fine-tune/Train Pruned Neural Network**
  - How should we improve performance of pruned models?





# Neural Network Pruning


- **Introduction to Pruning**
  - What is pruning?
  - How should we formulate pruning?
- **Determine the Pruning Granularity**
  - In what pattern should we prune the neural network?
- **Determine the Pruning Criterion**
  - What synapses/neurons should we prune?
- **Determine the Pruning Ratio**
  - What should target sparsity be for each layer?
- **Fine-tune/Train Pruned Neural Network**
  - How should we improve performance of pruned models?



Learning Both Weights and Connections for Efficient Neural Network [Han *et al.*, NeurIPS 2015]

# Neural Network Pruning


- **Introduction to Pruning**
  - What is pruning?
  - How should we formulate pruning?
- **Determine the Pruning Granularity**
  - In what pattern should we prune the neural network?
- **Determine the Pruning Criterion**
  - What synapses/neurons should we prune?
- **Determine the Pruning Ratio**
  - What should target sparsity be for each layer?
- **Fine-tune/Train Pruned Neural Network**
  - How should we improve performance of pruned models?




Learning Both Weights and Connections for Efficient Neural Network [Han *et al.*, NeurIPS 2015]

# Neural Network Pruning

- **Introduction to Pruning**
  - What is pruning?
  - How should we formulate pruning?
- **Determine the Pruning Granularity**
  - In what pattern should we prune the neural network?
- **Determine the Pruning Criterion**
  - What synapses/neurons should we prune?
- **Determine the Pruning Ratio**
  - What should target sparsity be for each layer?
- **Fine-tune/Train Pruned Neural Network**
  - How should we improve performance of pruned models?




Pruning  
.....➔



$$\begin{aligned} &\downarrow \\ &\arg \min_{\mathbf{W}_P} L(\mathbf{x}; \mathbf{W}_P) \\ &s.t. \|\mathbf{W}_P\|_0 \leq N \end{aligned}$$

# Neural Network Pruning

- **Introduction to Pruning**
  - What is pruning?
  - How should we formulate pruning?
- **Determine the Pruning Granularity**
  - In what pattern should we prune the neural network?
- **Determine the Pruning Criterion**
  - What synapses/neurons should we prune?
- **Determine the Pruning Ratio**
  - What should target sparsity be for each layer?
- **Fine-tune/Train Pruned Neural Network**
  - How should we improve performance of pruned models?






# Section 2: Pruning Granularity

**Pruning can be performed at different granularities, from structured to non-structured.**


# Pruning at Different Granularities



A simple example of 2D weight matrix



# Pruning at Different Granularities

A simple example of 2D weight matrix




 Preserved  
 Pruned

## Fine-grained/Unstructured

- More flexible pruning index choice
- Hard to accelerate (irregular)


# Pruning at Different Granularities

A simple example of 2D weight matrix



**Fine-grained/Unstructured**


- More flexible pruning index choice
- Hard to accelerate (irregular)



**Coarse-grained/Structured**

- Less flexible pruning index choice (a subset of the fine-grained case)
- Easy to accelerate (just a smaller matrix!)

■ Preserved  
□ Pruned





# Pruning at Different Granularities

## The case of convolutional layers

- The weights of convolutional layers have 4 dimensions  $[c_o, c_i, k_h, k_w]$ :
  - $c_i$ : input channels (or channels)
  - $c_o$ : output channels (or filters)
  - $k_h$ : kernel size height
  - $k_w$ : kernel size width
- The 4 dimensions give us more choices to select pruning granularities

# Pruning at Different Granularities


## The case of convolutional layers

- Some of the commonly used pruning granularities

■ Preserved

□ Pruned

### Notations




# Pruning at Different Granularities

## The case of convolutional layers


- Some of the commonly used pruning granularities

■ Preserved  
□ Pruned




Irregular

Regular




Fine-grained  
Pruning




Pattern-based  
Pruning


like Tetris :)



Vector-level  
Pruning



Kernel-level  
Pruning



Channel-level  
Pruning

Exploring the granularity of sparsity in convolutional neural networks [Mao *et al.*, CVPR-W]




# Pruning at Different Granularities

## The case of convolutional layers

- Some of the commonly used pruning granularities

■ Preserved  
□ Pruned




Pros?


Cons?

Irregular


Regular




Fine-grained  
Pruning




Pattern-based  
Pruning



Vector-level  
Pruning



Kernel-level  
Pruning




Channel-level  
Pruning

Exploring the granularity of sparsity in convolutional neural networks [Mao *et al.*, CVPR-W]

# Pruning at Different Granularities

## Let's look into some cases

- **Fine-grained Pruning** (the case we show before)
  - Flexible pruning indices



Learning Both Weights and Connections for Efficient Neural Network [Han *et al.*, NeurIPS 2015]

# Pruning at Different Granularities

## Let's look into some cases

- **Fine-grained Pruning** (the case we show before)
  - Flexible pruning indices
  - Usually larger compression ratio since we can flexibly find “redundant” weights (we will later discuss how we find them)

Neural Network	#Parameters		
	Before Pruning	After Pruning	Reduction
AlexNet	61 M	6.7 M	9 ×
VGG-16	138 M	10.3 M	12 ×
GoogleNet	7 M	2.0 M	3.5 ×
ResNet50	26 M	7.47 M	3.4 ×

# Pruning at Different Granularities

## Let's look into some cases

- **Fine-grained Pruning** (the case we show before)
  - Flexible pruning indices
  - Usually larger compression ratio since we can flexibly find “redundant” weights (we will later discuss how we find them)
  - Can deliver speed up on some custom hardware (e.g., EIE) but not GPU (easily)




# Pruning at Different Granularities

## The case of convolutional layers

- Some of the commonly used pruning granularities

■ Preserved  
□ Pruned




Pros?


Cons?

Irregular


Regular




Fine-grained  
Pruning




Pattern-based  
Pruning



Vector-level  
Pruning



Kernel-level  
Pruning




Channel-level  
Pruning

Exploring the granularity of sparsity in convolutional neural networks [Mao et al., CVPR-W]

# Pruning at Different Granularities

Let's look into some cases

- **Pattern-based Pruning: N:M sparsity**
  - N:M sparsity means that in each contiguous M elements, N of them is pruned




Dense Matrix


# Pruning at Different Granularities

Let's look into some cases

- **Pattern-based Pruning: N:M sparsity**
  - N:M sparsity means that in each contiguous M elements, N of them is pruned
  - A classic case is 2:4 sparsity (50% sparsity)



Dense Matrix




2:4 Sparse Matrix

# Pruning at Different Granularities

## Let's look into some cases

- **Pattern-based Pruning: N:M sparsity**

- N:M sparsity means that in each contiguous M elements, N of them is pruned
- A classic case is 2:4 sparsity (50% sparsity)
- It is supported by NVIDIA's Ampere GPU Architecture, which delivers up to 2x speed up



Accelerating Inference with Sparsity Using the NVIDIA Ampere Architecture and NVIDIA TensorRT



# Pruning at Different Granularities

## Let's look into some cases

- **Pattern-based Pruning: N:M sparsity**
  - N:M sparsity means that in each contiguous M elements, N of them is pruned
  - A classic case is 2:4 sparsity (50% sparsity)
  - It is supported by NVIDIA's Ampere GPU Architecture, which delivers ~2x speed up
  - Usually maintains accuracy (tested on varieties of tasks)

Network	Data Set	Metric	Dense FP16	Sparse FP16
ResNet-50	ImageNet	Top-1	76.1	76.2
ResNeXt-101_32x8d	ImageNet	Top-1	79.3	79.3
Xception	ImageNet	Top-1	79.2	79.2
SSD-RN50	COCO2017	bbAP	24.8	24.8
MaskRCNN-RN50	COCO2017	bbAP	37.9	37.9
FairSeq Transformer	EN-DE WMT'14	BLEU	28.2	28.5
BERT-Large	SQuAD v1.1	F1	91.9	91.9

Accelerating Inference with Sparsity Using the NVIDIA Ampere Architecture and NVIDIA TensorRT


# Pruning at Different Granularities

## The case of convolutional layers

- Some of the commonly used pruning granularities

■ Preserved

□ Pruned




Pros?


Cons?

Irregular


Regular




Fine-grained  
Pruning




Pattern-based  
Pruning



Vector-level  
Pruning



Kernel-level  
Pruning



Channel-level  
Pruning

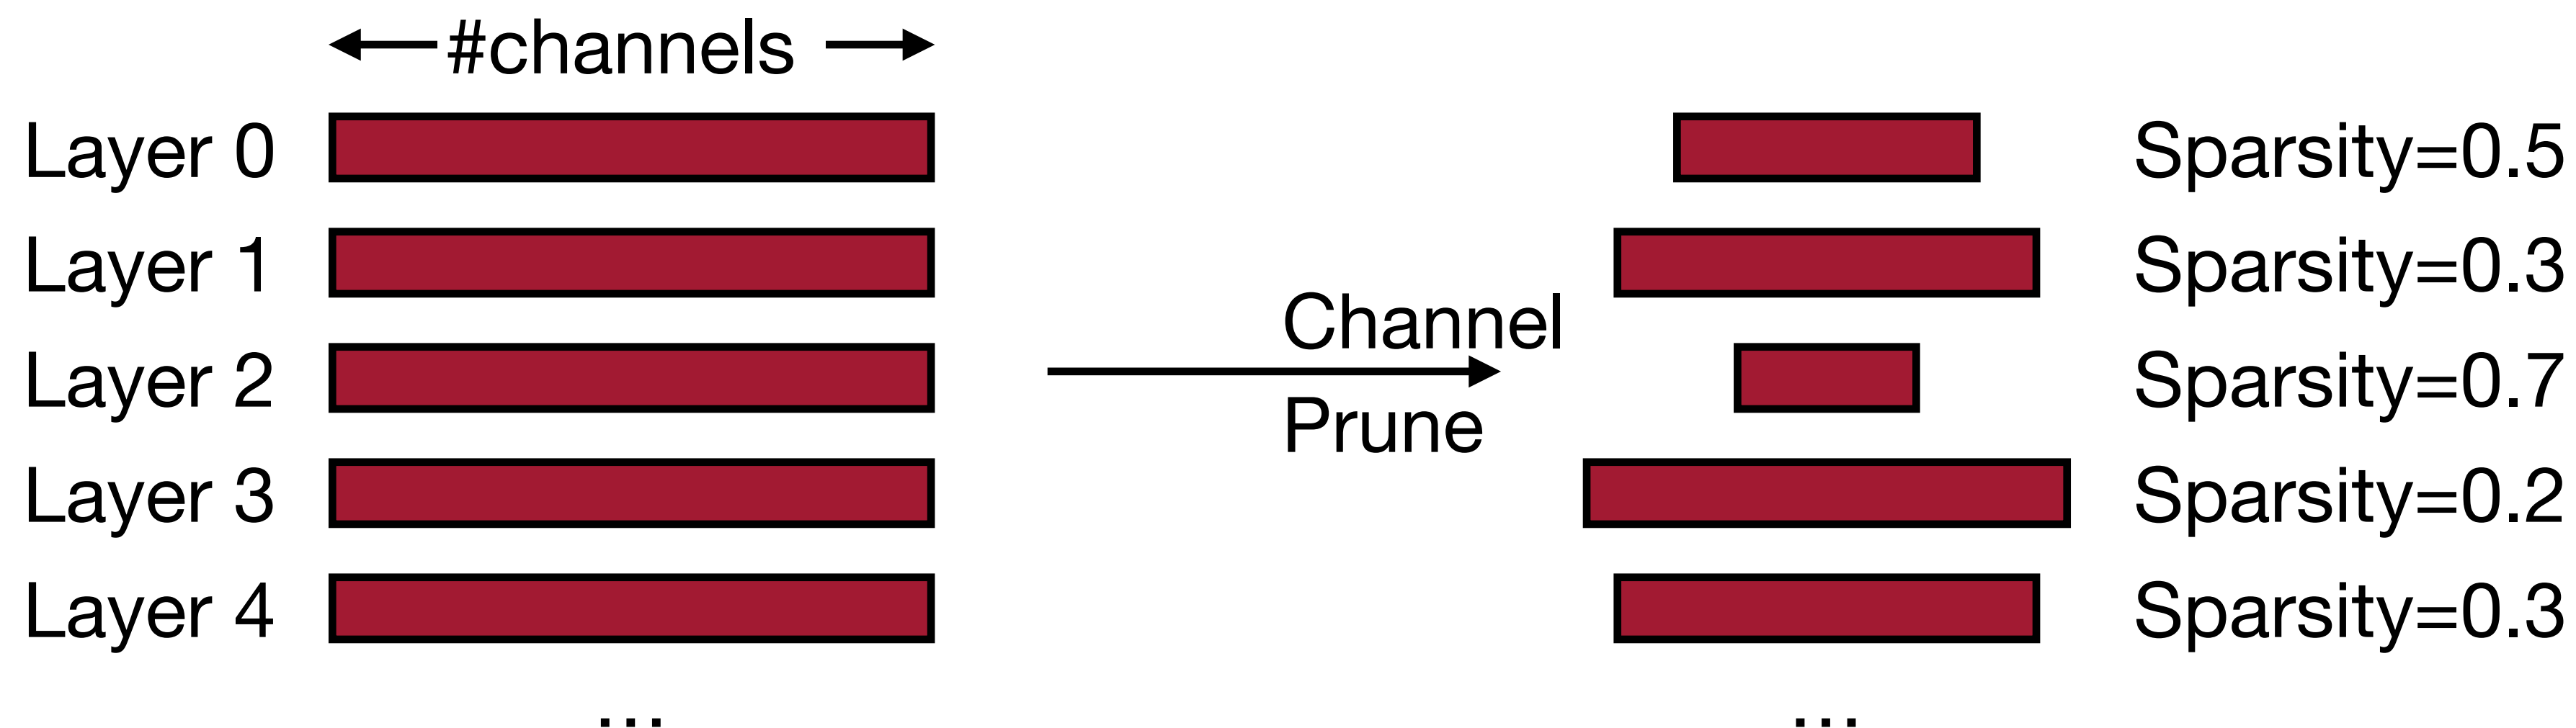
Exploring the granularity of sparsity in convolutional neural networks [Mao *et al.*, CVPR-W]

# Pruning at Different Granularities

## Let's look into some cases

- **Channel Pruning**

- Pro: Direct speed up due to reduced channel numbers (leading to an NN with smaller #channels)
- Con: smaller compression ratio








# Pruning at Different Granularities

## Let's look into some cases

- **Channel Pruning**

- Pro: Direct speed up due to reduced channel numbers (leading to an NN with smaller #channels)
- Con: smaller compression ratio



*We will later discuss how to find sparsity ratios*

Sparsity=0.3   
Sparsity=0.3   
Sparsity=0.3   
Sparsity=0.3   
Sparsity=0.3 

...

Uniform Shrink

<

 Sparsity=0.5  
 Sparsity=0.3  
 Sparsity=0.7  
 Sparsity=0.2  
 Sparsity=0.3

...

Channel Prune





# Pruning at Different Granularities

Let's look into some cases

- **Channel Pruning**


- Pro: Direct speed up due to reduced channel numbers (leading to an NN with smaller #channels)
- Con: smaller compression ratio



AMC: Automl for Model Compression and Acceleration on Mobile Devices [He *et al.*, ECCV 2018]

# Neural Network Pruning

- **Introduction to Pruning**
  - What is pruning?
  - How should we formulate pruning?
- **Determine the Pruning Granularity**
  - In what pattern should we prune the neural network?
- **Determine the Pruning Criterion**
  - What synapses/neurons should we prune?
- **Determine the Pruning Ratio**
  - What should target sparsity be for each layer?
- **Fine-tune/Train Pruned Neural Network**
  - How should we improve performance of pruned models?




Learning Both Weights and Connections for Efficient Neural Network [Han *et al.*, NeurIPS 2015]

# Section 3: Pruning Criterion

**What synapses and neurons should we prune?**

# Selection of Synapses to Prune

- When removing parameters from a neural network model,
  - ***the less important*** the parameters being removed are,
  - the better the performance of pruned neural network is.



## Example

$$f(\cdot) = \text{ReLU}(\cdot), \quad W = [10, -8, 0.1]$$

$$\Rightarrow y = \text{ReLU}(10x_0 - 8x_1 + 0.1x_2)$$

- If one weight will be removed, which one?



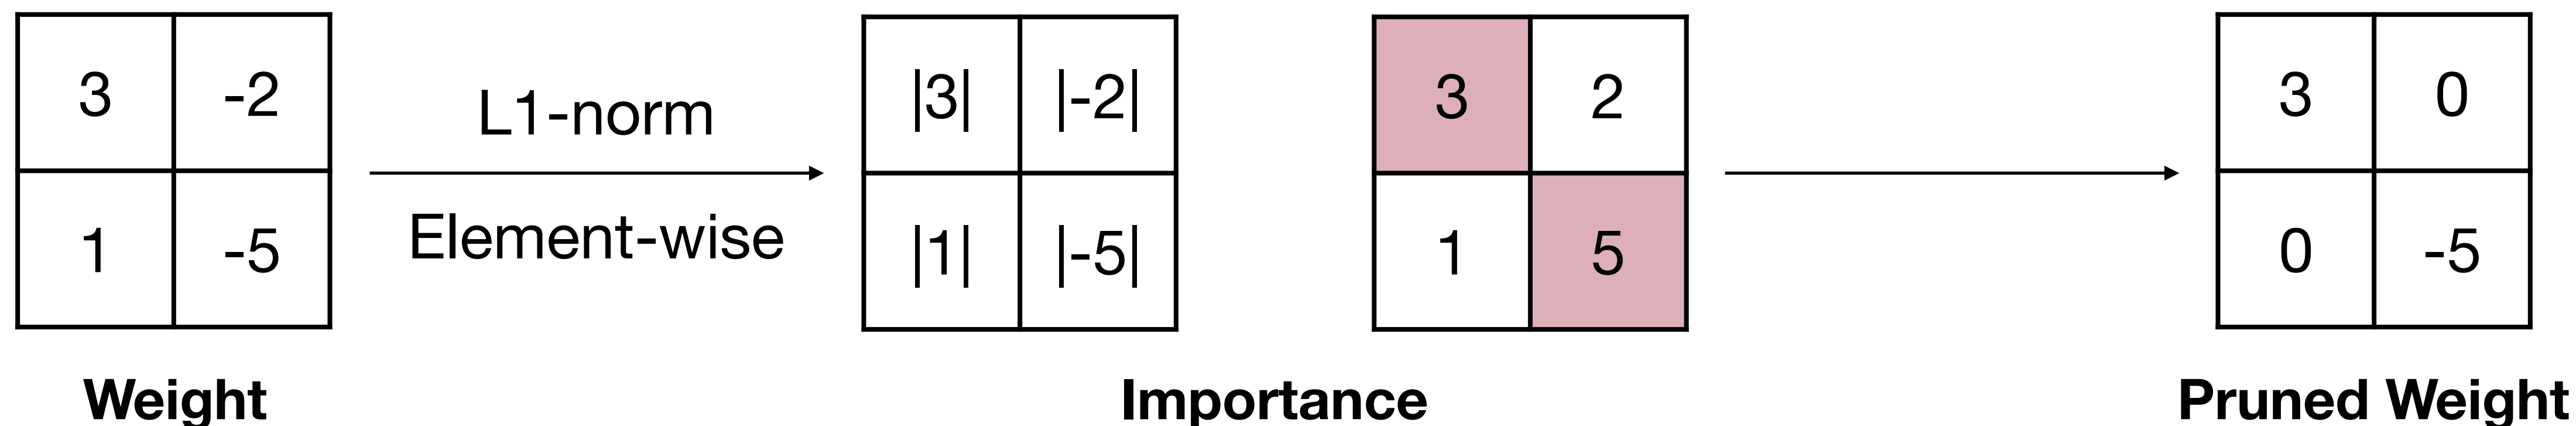
# Magnitude-based Pruning

## A heuristic pruning criterion

- Magnitude-based pruning considers weights with ***larger absolute values*** are more important than other weights.
  - For element-wise pruning,

$$\text{Importance} = |W|$$

- **Example**



Learning Both Weights and Connections for Efficient Neural Network [Han *et al.*, NeurIPS 2015]

# Magnitude-based Pruning

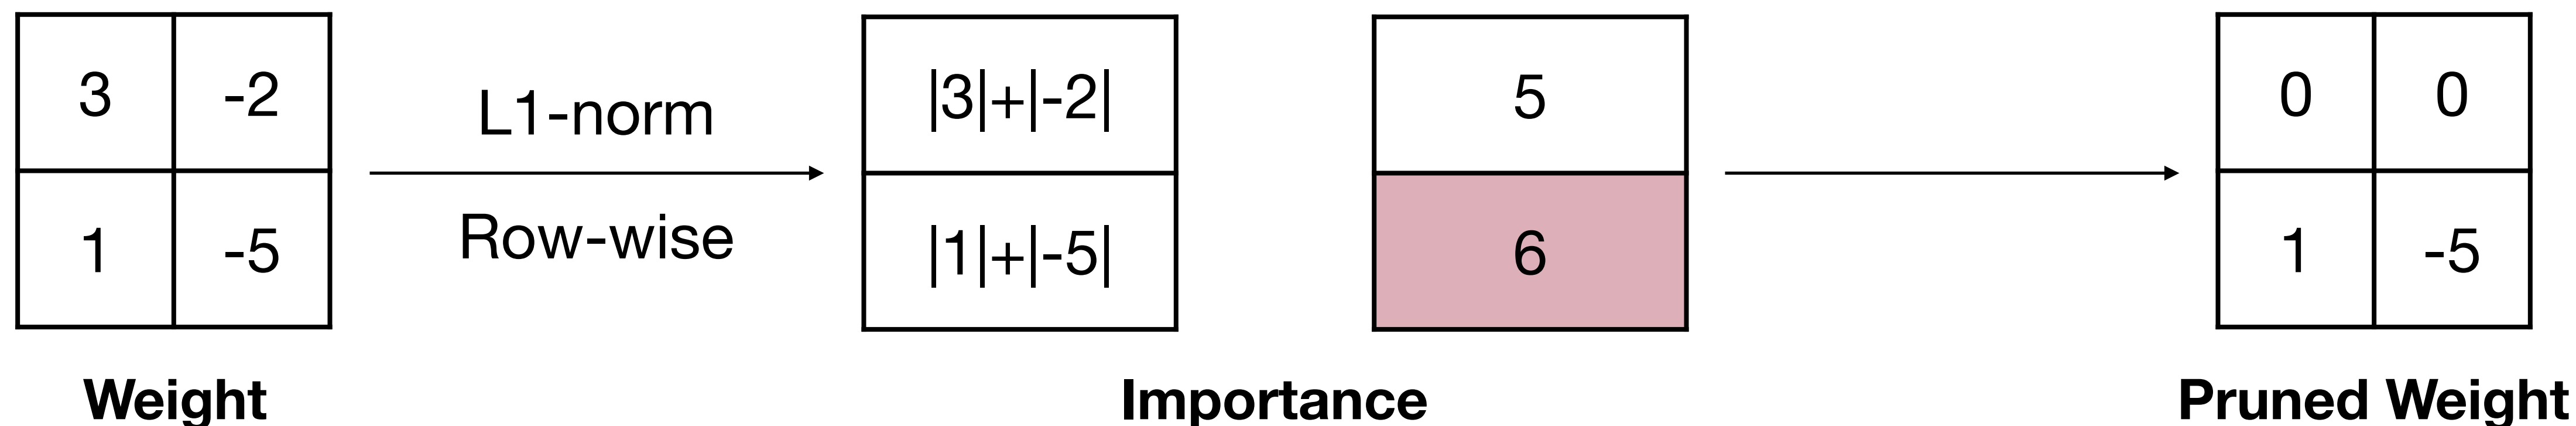
## A heuristic pruning criterion

- Magnitude-based pruning considers weights with ***larger absolute values*** are more important than other weights.

- For row-wise pruning, the L1-norm magnitude can be defined as,

$$Importance = \sum_{i \in S} |w_i|, \text{ where } \mathbf{W}^{(S)} \text{ is the structural set } S \text{ of parameters } \mathbf{W}$$

- **Example**



Learning Both Weights and Connections for Efficient Neural Network [Han *et al.*, NeurIPS 2015]

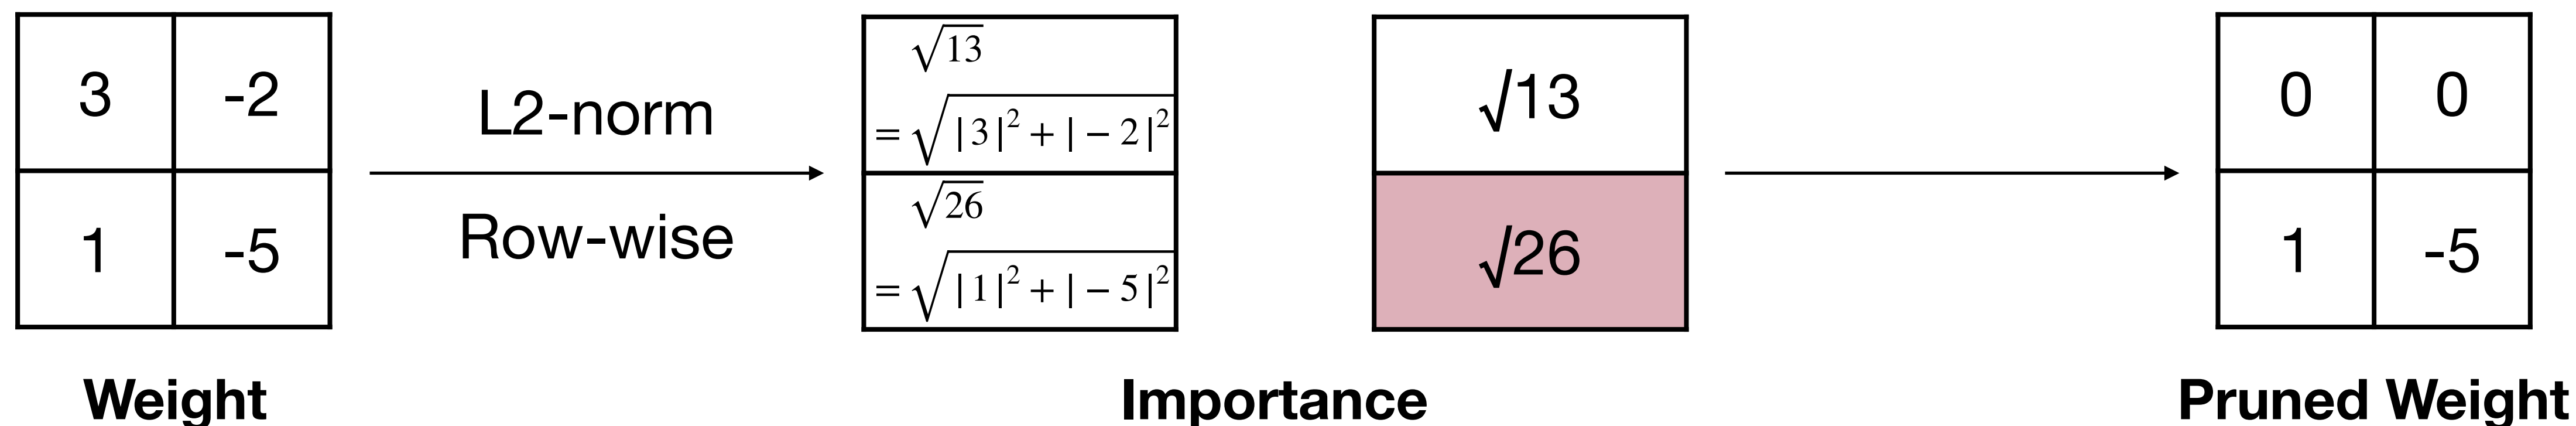
# Magnitude-based Pruning

## A heuristic pruning criterion

- Magnitude-based pruning considers weights with ***larger absolute values*** are more important than other weights.
- For row-wise pruning, the L2-norm magnitude can be defined as,

$$Importance = \sqrt{\sum_{i \in S} |w_i|^2}, \text{ where } \mathbf{W}^{(S)} \text{ is the structural set } S \text{ of parameters } \mathbf{W}$$

- **Example**



Learning Both Weights and Connections for Efficient Neural Network [Han *et al.*, NeurIPS 2015]

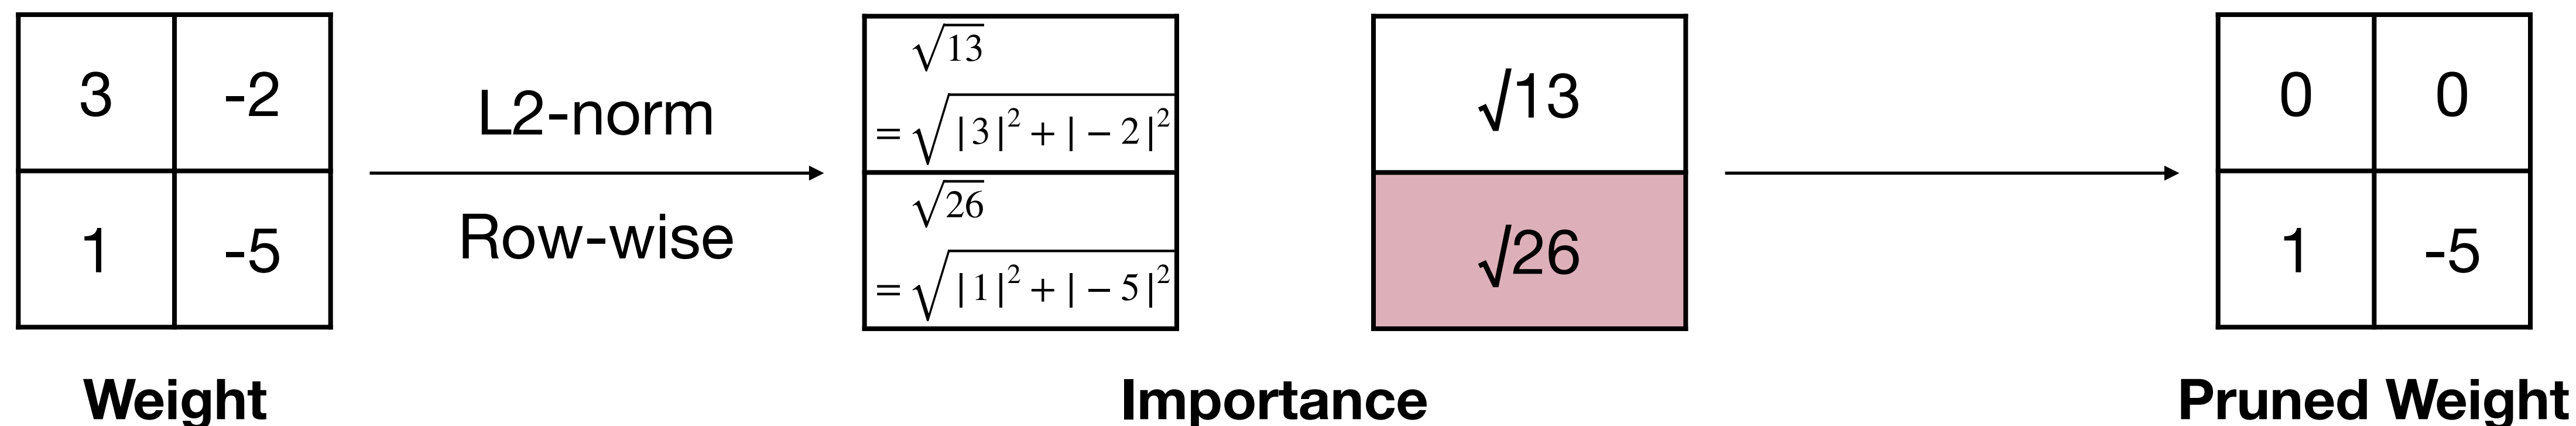
# Magnitude-based Pruning

## A heuristic pruning criterion

- Magnitude-based pruning considers weights with ***larger absolute values*** are more important than other weights.
- Magnitude is also known as  $L_p$ -norm defined as,

$$\|\mathbf{W}^{(S)}\|_p = \left( \sum_{i \in S} |w_i|^p \right)^{\frac{1}{p}}, \text{ where } \mathbf{W}^{(S)} \text{ is a structural set of parameters}$$

- **Example**

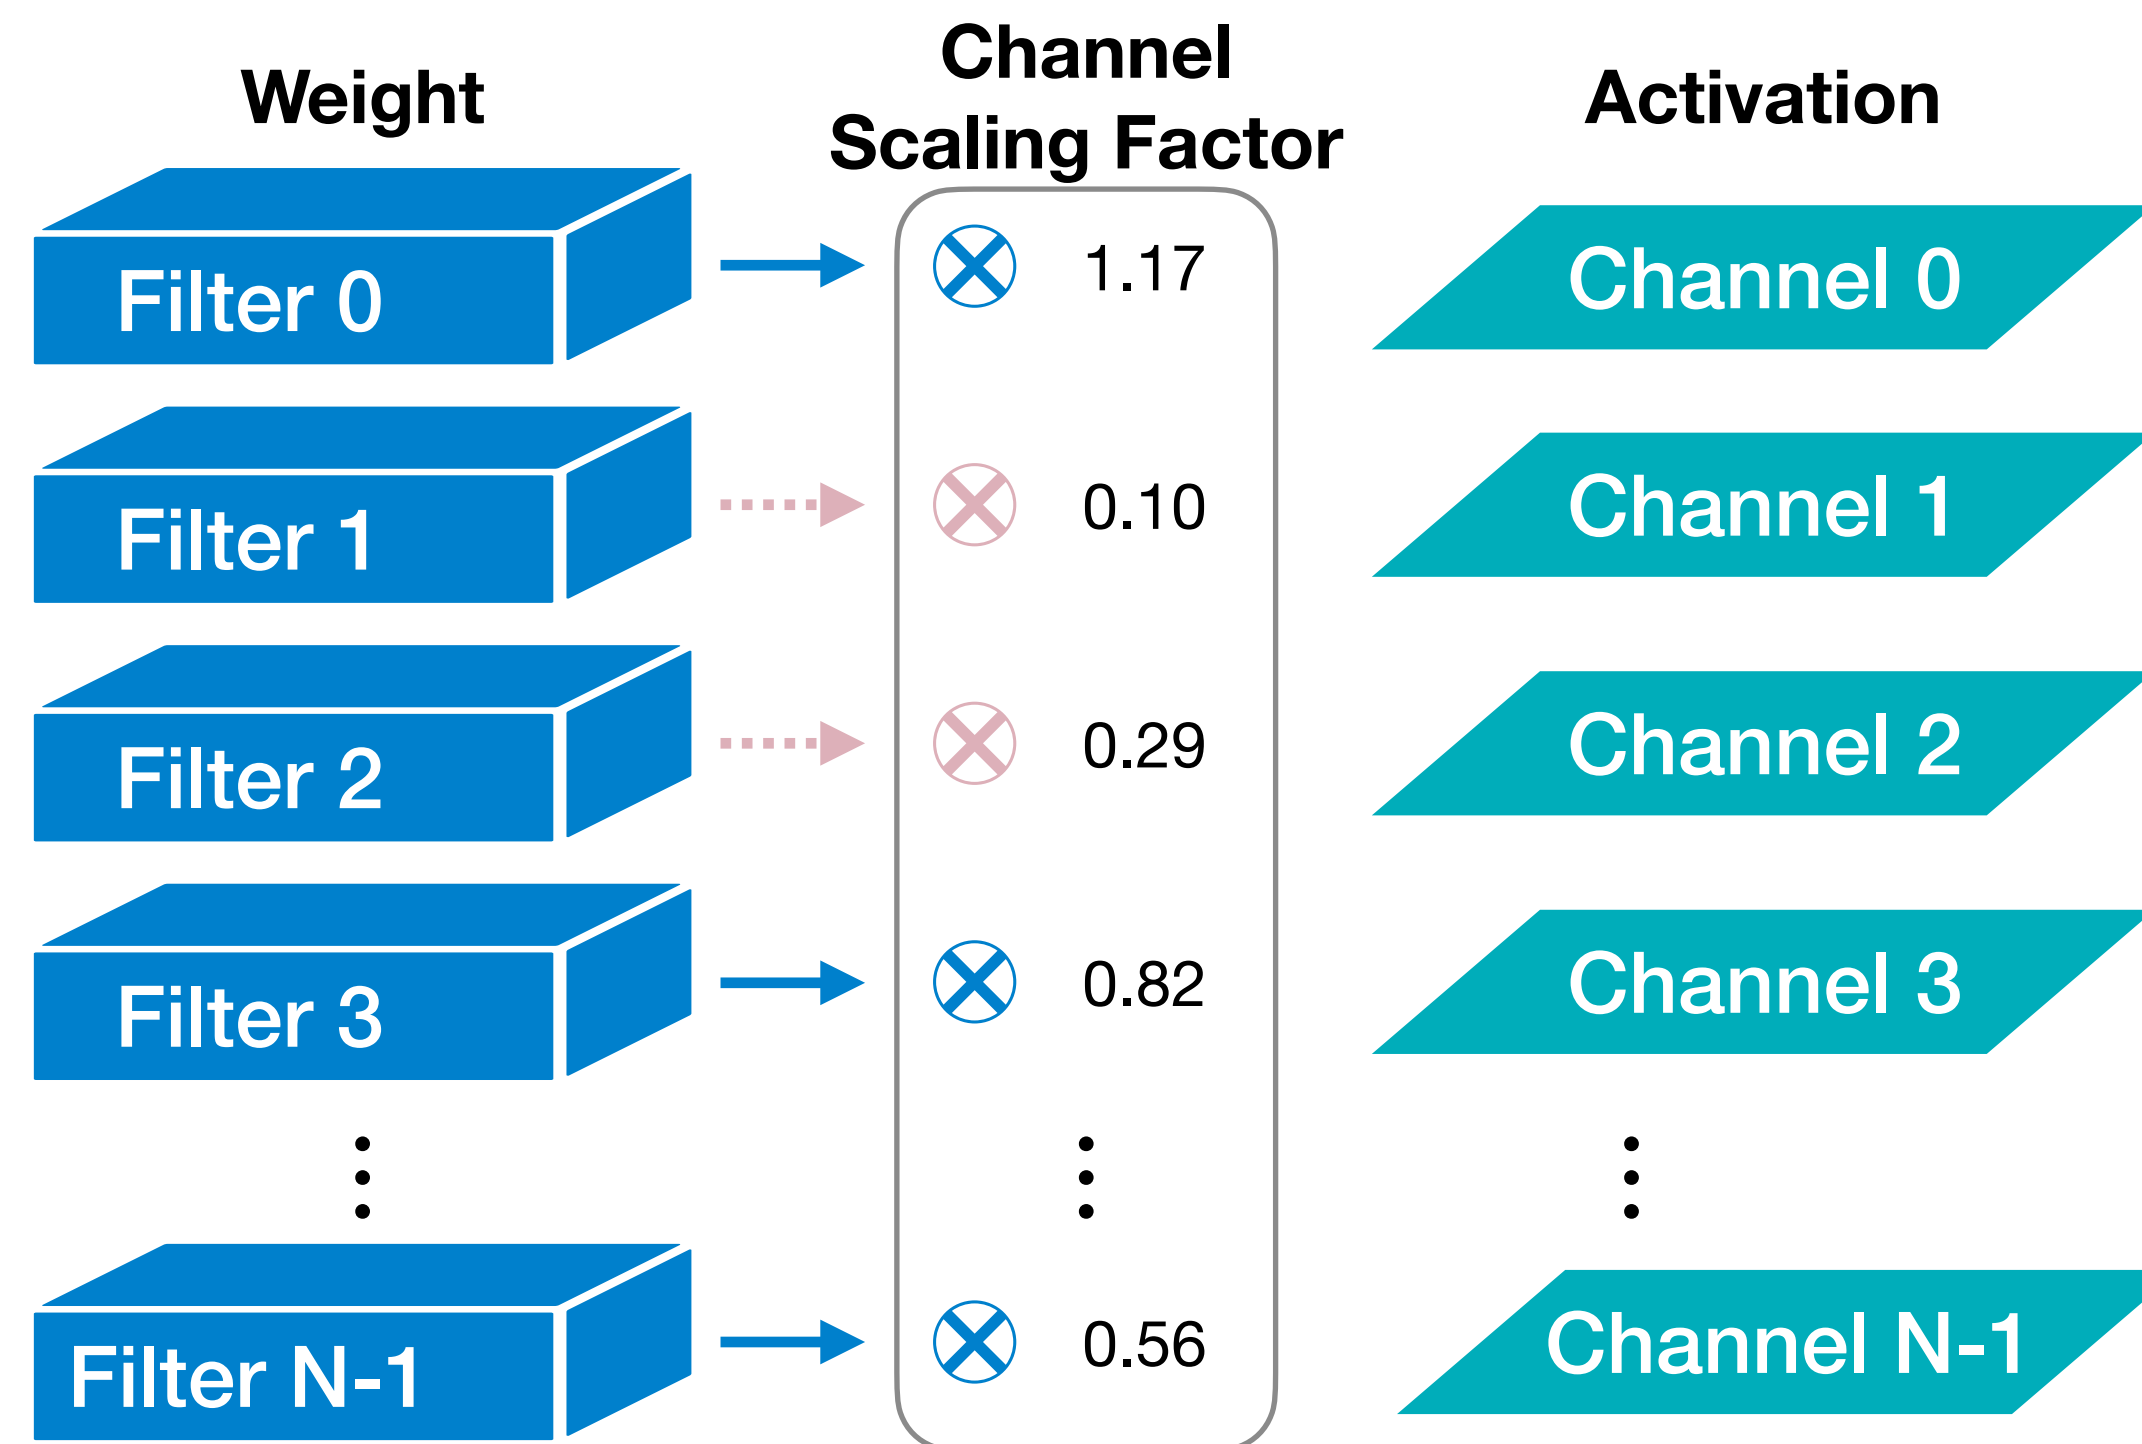




# Scaling-based Pruning

## Pruning criterion for filter pruning

- A scaling factor is associated with each filter (*i.e.*, output channel) in convolutional layers
  - The scaling factor is multiplied to the output of that channel
  - The scaling factors are trainable parameters

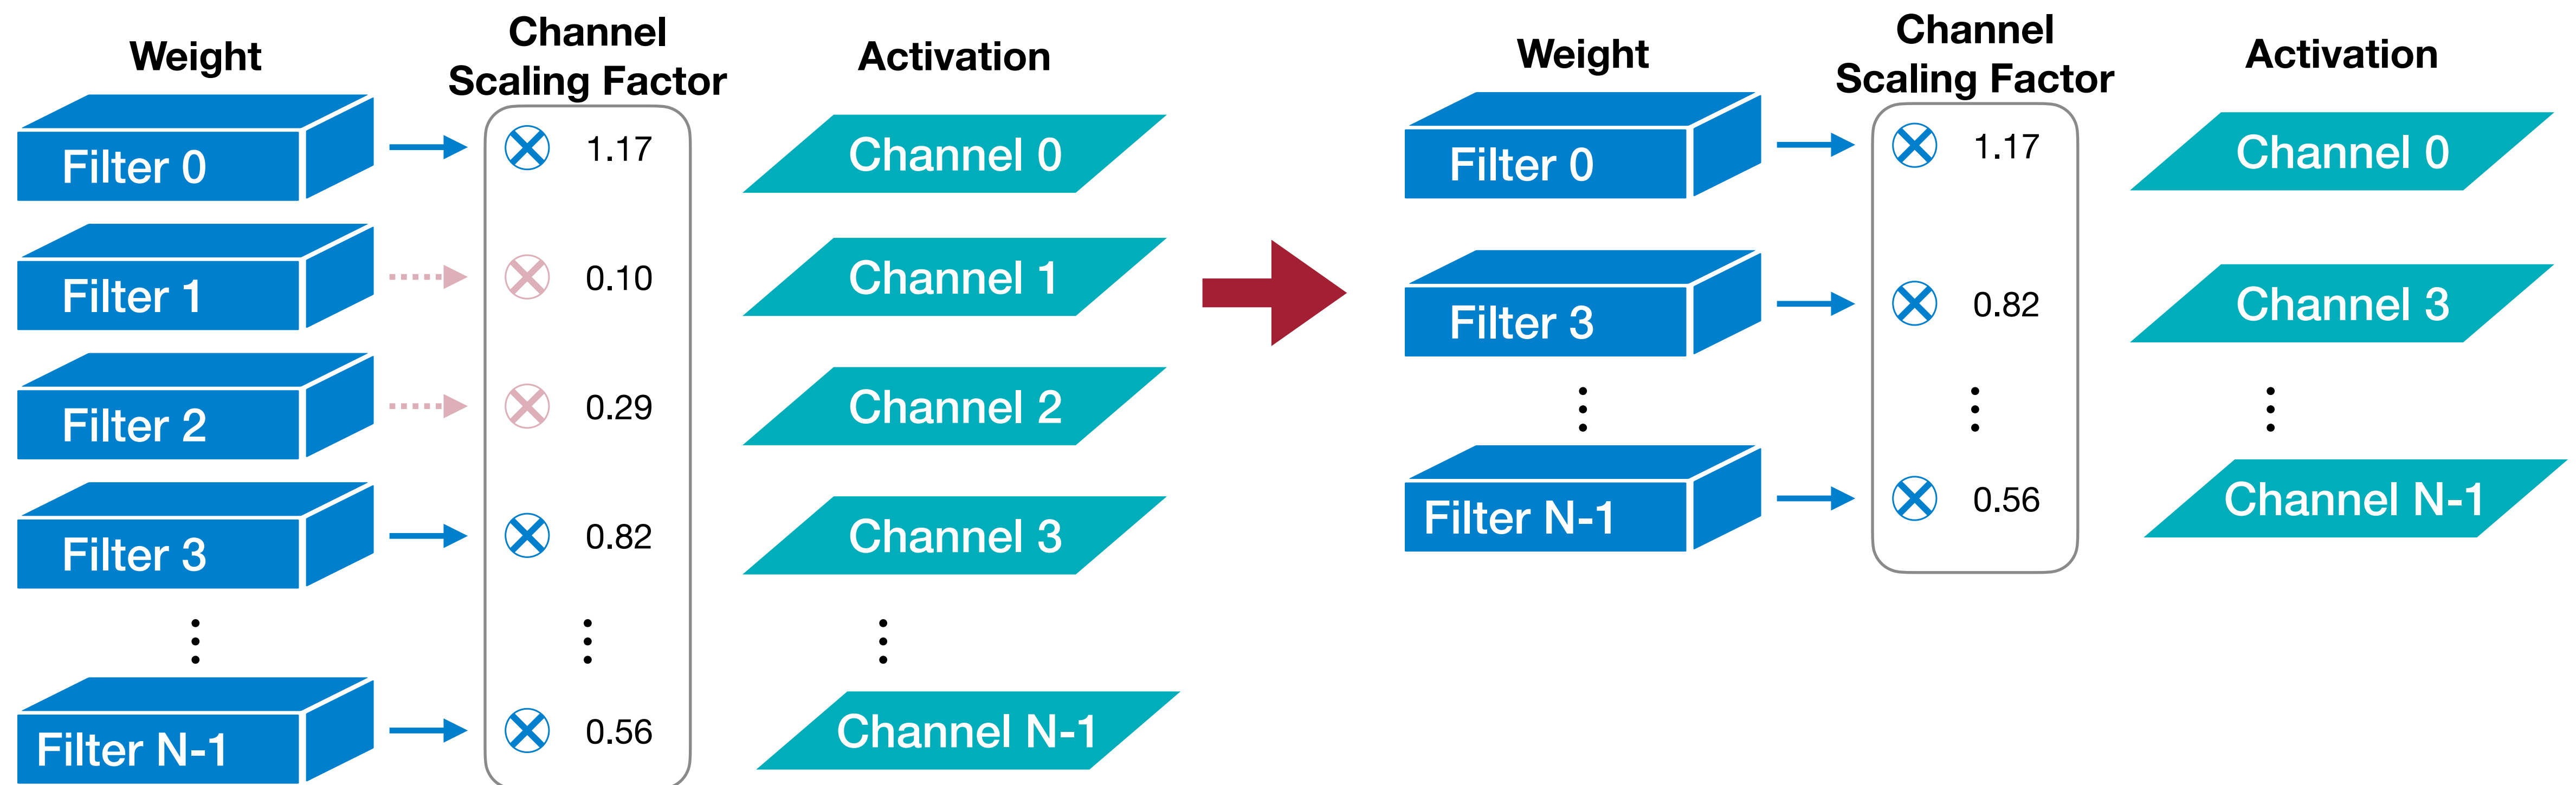


Learning Efficient Convolutional Networks through Network Slimming [Liu *et al.*, ICCV 2017]

# Scaling-based Pruning

## Pruning criterion for filter pruning

- A scaling factor is associated with each filter (*i.e.*, output channel) in convolutional layers
  - The scaling factor is multiplied to the output of that channel
  - The scaling factors are trainable parameters
- The filters/output channels with small scaling factor magnitude will be pruned



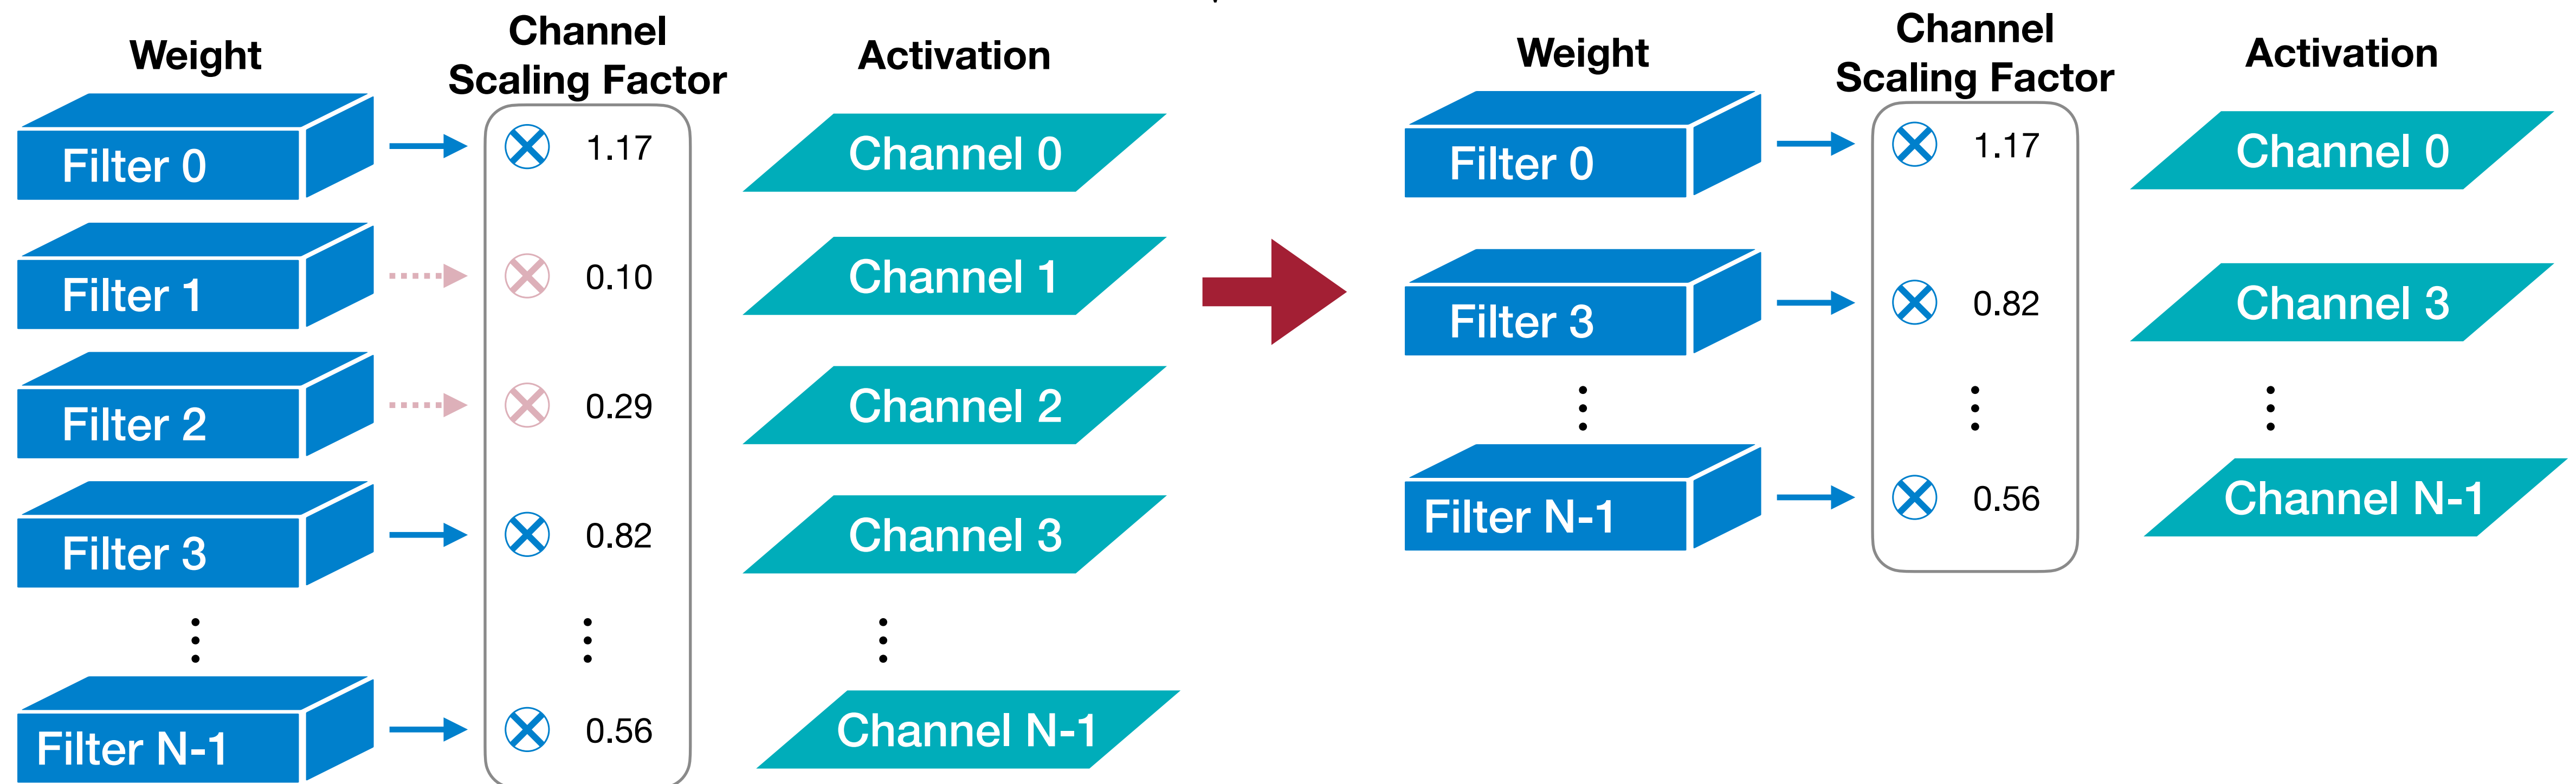
Learning Efficient Convolutional Networks through Network Slimming [Liu *et al.*, ICCV 2017]

# Scaling-based Pruning

## Pruning criterion for filter pruning

- A scaling factor is associated with each filter (*i.e.*, output channel) in convolutional layers
- The scaling factors can be reused from batch normalization layer

$$\mathbf{z}_o = \gamma \frac{\mathbf{z}_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} + \beta$$



Learning Efficient Convolutional Networks through Network Slimming [Liu *et al.*, ICCV 2017]

# Second-Order-based Pruning

**Minimize the error on loss function introduced by pruning synapses**

- The induced error can be approximated by a Taylor series.

$$\delta L = L(\mathbf{x}; \mathbf{W}) - L(\mathbf{x}; \mathbf{W}_P = \mathbf{W} - \delta \mathbf{W}) = \sum_i g_i \delta w_i + \frac{1}{2} \sum_i h_{ii} \delta w_i^2 + \frac{1}{2} \sum_{i \neq j} h_{ij} \delta w_i \delta w_j + O(\|\delta \mathbf{W}\|^3)$$

where

$$g_i = \frac{\partial L}{\partial w_i}, h_{i,j} = \frac{\partial^2 L}{\partial w_i \partial w_j}$$

- Optimal Brain Damage assumes that



# Second-Order-based Pruning

**Minimize the error on loss function introduced by pruning synapses**

- The induced error can be approximated by a Taylor series.

$$\delta L = L(\mathbf{x}; \mathbf{W}) - L(\mathbf{x}; \mathbf{W}_P = \mathbf{W} - \delta \mathbf{W}) = \sum_i g_i \delta w_i + \frac{1}{2} \sum_i h_{ii} \delta w_i^2 + \frac{1}{2} \sum_{i \neq j} h_{ij} \delta w_i \delta w_j + O(\|\delta \mathbf{W}\|^3)$$

where

$$g_i = \frac{\partial L}{\partial w_i}, h_{i,j} = \frac{\partial^2 L}{\partial w_i \partial w_j}$$

- Optimal Brain Damage assumes that
  - The objective function  $L$  is nearly quadratic: the last term is neglected

# Second-Order-based Pruning

**Minimize the error on loss function introduced by pruning synapses**

- The induced error can be approximated by a Taylor series.

$$\delta L = L(\mathbf{x}; \mathbf{W}) - L(\mathbf{x}; \mathbf{W}_P = \mathbf{W} - \delta \mathbf{W}) = \sum_i g_i \delta w_i + \frac{1}{2} \sum_i h_{ii} \delta w_i^2 + \frac{1}{2} \sum_{i \neq j} h_{ij} \delta w_i \delta w_j + O(\|\delta \mathbf{W}\|^3)$$

where

$$g_i = \frac{\partial L}{\partial w_i}, h_{i,j} = \frac{\partial^2 L}{\partial w_i \partial w_j}$$

- Optimal Brain Damage assumes that
  - The objective function  $L$  is nearly quadratic: the last term is neglected
  - The neural network training has converged: first-order terms are neglected

# Second-Order-based Pruning

Minimize the error on loss function introduced by pruning synapses

- The induced error can be approximated by a Taylor series.

$$\delta L = L(\mathbf{x}; \mathbf{W}) - L(\mathbf{x}; \mathbf{W}_P = \mathbf{W} - \delta \mathbf{W}) = \sum_i g_i \delta w_i + \frac{1}{2} \sum_i h_{ii} \delta w_i^2 + \frac{1}{2} \sum_{i \neq j} h_{ij} \delta w_i \delta w_j + O(\|\delta \mathbf{W}\|^3)$$

where

$$g_i = \frac{\partial L}{\partial w_i}, h_{i,j} = \frac{\partial^2 L}{\partial w_i \partial w_j}$$

- Optimal Brain Damage assumes that
  - The objective function  $L$  is nearly quadratic: the last term is neglected
  - The neural network training has converged: first-order terms are neglected
  - The error caused by deleting each parameter is independent: cross terms are neglected

# Second-Order-based Pruning

Minimize the error on loss function introduced by pruning synapses

- The induced error can be approximated by a Taylor series.

$$\delta L = L(\mathbf{x}; \mathbf{W}) - L(\mathbf{x}; \mathbf{W}_P = \mathbf{W} - \delta \mathbf{W}) = \sum_i g_i \delta w_i + \frac{1}{2} \sum_i h_{ii} \delta w_i^2 + \frac{1}{2} \sum_{i \neq j} h_{ij} \delta w_i \delta w_j + O(\|\delta \mathbf{W}\|^3)$$

where

$$g_i = \frac{\partial L}{\partial w_i}, h_{i,j} = \frac{\partial^2 L}{\partial w_i \partial w_j}$$

- Optimal Brain Damage assumes that
  - The objective function  $L$  is nearly quadratic: the last term is neglected
  - The neural network training has converged: first-order terms are neglected
  - The error caused by deleting each parameter is independent: cross terms are neglected

$$\delta L_i = L(\mathbf{x}; \mathbf{W}) - L(\mathbf{x}; \mathbf{W}_P | w_i = 0) \approx \frac{1}{2} h_{ii} w_i^2$$

Optimal Brain Damage [LeCun et al., NeurIPS 1989]



# Second-Order-based Pruning

Minimize the error on loss function introduced by pruning synapses

- Optimal Brain Damage assumes that
  - The objective function  $L$  is nearly quadratic
  - The neural network training has converged
  - The error caused by deleting each parameter is independent

$$\delta L_i = L(\mathbf{x}; \mathbf{W}) - L(\mathbf{x}; \mathbf{W}_P | w_i = 0) \approx \frac{1}{2} h_{ii} w_i^2, \quad \text{where } h_{ii} = \frac{\partial^2 L}{\partial w_i \partial w_i}$$

- The synapses with smaller induced error  $|\delta L_i|$  will be removed; that is to say,

$$importance_{w_i} = |\delta L_i| = \frac{1}{2} h_{ii} w_i^2$$

\*  $h_{ii}$  is non-negative

**Hessian Matrix  $H$  is difficult to compute.**


Optimal Brain Damage [LeCun *et al.*, NeurIPS 1989]

# Selection of Neurons to Prune


- When removing neurons from a neural network model,
  - ***the less useful*** the neurons being removed are,
  - the better the performance of pruned neural network is.

Neuron pruning is coarse-grained weight pruning



Weight Matrix



Neuron Pruning  
in Linear Layer



Channel Pruning  
in Convolution Layer



# Percentage-of-Zero-Based Pruning

- ReLU activation will generate zeros in the output activation.

Output Activations	Height = 4	Width = 4				Channel = 3	Batch = 2	Width = 4				Channel = 3
		0	0.1	0.5	1			0.1	0.5	0	0	
		1.2	0.6	0.3	0.2			0.2	0.3	0	1	
		0	0.5	0	0.3			0.1	0	0	0.5	
		0.2	0	0	0.8			0.1	0.6	0.7	0.1	
	Height = 4	0.5	0	0.2	0.1	Channel = 3	Batch = 2	0.1	0.5	0	0	
		0	0.2	1.2	0			0	0.8	0	1	
		1.2	0	0.2	0.3			0.1	0	0.1	1.0	
		0.2	0.4	0	0			0.2	0	1.0	0	
		0	0.8	0.1	0			0.2	0	0.3	0	

# Percentage-of-Zero-Based Pruning

- ReLU activation will generate zeros in the output activation.
- Similar to magnitude of weights, the Average Percentage of Zero activations (APoZ) can be exploited to measure the importance of the neurons.

Output Activations		Width = 4											
		Height = 4											
						Channel = 3				Batch = 2			
										Channel = 3			

Average Percentage of Zeros (APoZ)

$$= \frac{5 + 6}{2 \cdot 4 \cdot 4} = \frac{11}{32}$$

Channel 0

$$= \frac{5 + 7}{2 \cdot 4 \cdot 4} = \frac{12}{32}$$

Channel 1

$$= \frac{6 + 8}{2 \cdot 4 \cdot 4} = \frac{14}{32}$$

Channel 2






# Regression-based Pruning

**Minimize reconstruction error of the corresponding layer's outputs**



- Instead of considering the pruning error of the objective function  $L(\mathbf{x}; \mathbf{W})$ , regression-based pruning minimizes the reconstruction error of the corresponding layer's outputs.



# Regression-based Pruning

Minimize reconstruction error of the corresponding layer's outputs


- Instead of considering the pruning error of the objective function  $L(\mathbf{x}; \mathbf{W})$ , regression-based pruning minimizes the reconstruction error of the corresponding layer's outputs.



# Regression-based Pruning

## Minimize reconstruction error of the corresponding layer's outputs

- Instead of considering the pruning error of the objective function  $L(\mathbf{x}; \mathbf{W})$ , regression-based pruning minimizes the reconstruction error of the corresponding layer's outputs.




# Regression-based Pruning

Minimize reconstruction error of the corresponding layer's outputs

- Let

$$\mathbf{Z} = \mathbf{X}\mathbf{W}^T = \sum_{c=0}^{c_i-1} \mathbf{X}_c \mathbf{W}_c^T$$





# Regression-based Pruning

Minimize reconstruction error of the corresponding layer's outputs

- Let

$$\mathbf{Z} = \mathbf{X}\mathbf{W}^T = \sum_{c=0}^{c_i-1} \mathbf{X}_c \mathbf{W}_c^T$$


- The problem can be formulate as

$$\arg \min_{\mathbf{W}, \beta} \|\mathbf{Z} - \hat{\mathbf{Z}}\|_F^2 = \|\mathbf{Z} - \sum_{c=0}^{c_i-1} \beta_c \mathbf{X}_c \mathbf{W}_c^T\|_F^2$$

subject to

$$\|\beta\|_0 \leq N_c$$

- $\beta$  is coefficient vector of length  $c_i$  for channel selection.  $\beta_c = 0$  means channel  $c$  is pruned.
- $N_c$  is the number of nonzero channels.



# Regression-based Pruning

Minimize reconstruction error of the corresponding layer's outputs

- Let

$$\mathbf{Z} = \mathbf{X}\mathbf{W}^T = \sum_{c=0}^{c_i-1} \mathbf{X}_c \mathbf{W}_c^T$$


- The problem can be formulate as

$$\arg \min_{\mathbf{W}, \beta} \|\mathbf{Z} - \hat{\mathbf{Z}}\|_F^2 = \|\mathbf{Z} - \sum_{c=0}^{c_i-1} \beta_c \mathbf{X}_c \mathbf{W}_c^T\|_F^2$$

subject to

$$\|\beta\|_0 \leq N_c$$

- $\beta$  is coefficient vector of length  $c_i$  for channel selection.  $\beta_c = 0$  means channel  $c$  is pruned.
- $N_c$  is the number of nonzero channels.
- Solve the problem by:
  - Fix  $\mathbf{W}$ , solve  $\beta$  for channel selection
  - Fix  $\beta$ , solve  $\mathbf{W}$  to minimize reconstruction error




Channel Pruning for Accelerating Very Deep Neural Networks [He et al., ICCV 2017]

# Summary of Today's Lecture

## Pruning Demo

In this lecture, we introduced:

- What is pruning
- Granularities of pruning
- Criteria to select weights to prune
- **We will cover in the next lecture:**
  - How to find pruning ratio for each layer
  - How to train/fine-tune the pruned layer
  - Automated ways to find pruning ratios
  - Lottery ticket hypothesis
  - System support for different granularities



# References

1. Model Compression and Hardware Acceleration for Neural Networks: A Comprehensive Survey [Deng et al., IEEE 2020]
2. Computing's Energy Problem (and What We Can Do About it) [Horowitz, M., IEEE ISSCC 2014]
3. Optimal Brain Damage [LeCun et al., NeurIPS 1989]
4. Learning Both Weights and Connections for Efficient Neural Network [Han et al., NeurIPS 2015]
5. Efficient Methods and Hardware for Deep Learning [Han S., Stanford University]
6. Peter Huttenlocher (1931–2013) [Walsh, C. A., Nature 2013]
7. Exploring the granularity of sparsity in convolutional neural networks [Mao et al., CVPR-W]
8. Accelerating Inference with Sparsity Using the NVIDIA Ampere Architecture and NVIDIA TensorRT
9. AMC: Automl for Model Compression and Acceleration on Mobile Devices [He et al., ECCV 2018]
10. Learning Structured Sparsity in Deep Neural Networks [Wen et al., NeurIPS 2016]
11. Learning Efficient Convolutional Networks through Network Slimming [Liu et al., ICCV 2017]
12. Pruning Convolutional Filters with First Order Taylor Series Ranking [Wang M.]
13. Importance Estimation for Neural Network Pruning [Molchanov et al., CVPR 2019]
14. Network Trimming: A Data-Driven Neuron Pruning Approach towards Efficient Deep Architectures [Hu et al., ArXiv 2017]
15. Pruning Convolutional Neural Networks for Resource Efficient Inference [Molchanov et al., ICLR 2017]
16. Channel Pruning for Accelerating Very Deep Neural Networks [He et al., ICCV 2017]
17. ThiNet: A Filter Level Pruning Method for Deep Neural Network Compression [Luo et al., ICCV 2017]
18. SparseGPT: Massive Language Models Can be Accurately Pruned in One-Shot [Elias Frantar, Dan Alistarh, ArXiv 2023]